

Type-safe Unions in C++14

Toby Allsopp
toby@mi6.gen.nz

Auckland C++ Meetup
19 July 2016

Introduction

Motivating
Example

Discriminated
Unions

Plain old `union`
`union + tag`

Inline visitors

Multi-
visitation

Generic
variant

Questions

This talk

Type-safe
Unions in
C++14

Toby Allsopp

Introduction

Motivating
Example

Discriminated
Unions

Plain old `union`
`union + tag`

Inline visitors

Multi-
visitation

Generic
variant

Questions

Overview

Introduction

Motivating Example

Discriminated Unions

Plain old **union**
union + tag

Inline visitors

Multi-visitation

Generic variant

Questions

Type-safe
Unions in
C++14

Toby Allsopp

Introduction

Motivating
Example

Discriminated
Unions

Plain old **union**
union + tag


Inline visitors

Multi-
visitation

Generic
variant

Questions

A robot



► States

```
struct off {};  
struct idle {};  
struct turning { float target; };
```

► Events

```
struct turn_on {};  
struct turn_off {};  
struct start_turning { float target; };  
struct heading_changed { float heading; };  
struct reset { std::string reason; };
```

- ▶ We want to write a state transition function for the robot

```
state transition(const state& s, const event& e);
```

- ▶ But what types do we use for state and event?
- ▶ We need something that can hold **either** an off, an idle or a turning (and similarly for the events)
- ▶ Something like a **union**
- ▶ But better!

variant

- ▶ Boost has `boost::variant`
- ▶ C++17 will have `std::variant`
- ▶ We're going to build our own

```
using state = variant<off, idle, turning>;  
using event =  
    variant<turn_on, turn_off, start_turning, reset,  
            heading_changed>;
```

Indiscriminate **union**

```
union u {  
    int i;  
    double d;  
}  
u x;  
x.i = 3;  
double d = x.d; // UNDEFINED BEHAVIOUR
```

- ▶ Can't tell what was last stored but you'd better know!
- ▶ Can be used for bit twiddling on specific platforms
- ▶ No guarantees from the standard

More **union** gotchas

```
union u {  
    std::string s;  
    std::vector<int> v;  
} x;  
x.s = "Hello"; // KABOOM!  
new (&x.s) std::string("Hello"); // OK  
x.s = "Goodbye"; // OK  
u y = x; // NOPE - copy constructor is deleted  
new (&x.v) std::vector<int>{1, 2, 3}; // LEAK!  
x.s.~std::string(); // DO THIS FIRST
```

Introduction

Motivating
Example

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Generic
variant

Questions

union + class + enum

```
class svu {  
    private:  
        enum { STRING, VECTOR } tag;  
        union {  
            std::string s;  
            std::vector<int> v;  
        };  
};
```

- ▶ Because it's so unsafe, let's bundle it up into a class and only expose safe operations.
- ▶ We need to keep track of which union member we have initialized - we call this a `tag`.

Construction

```
private:
```

```
void construct(const std::string& _s) {  
    tag = STRING;  
    new (&s) std::string(_s);  
}  
void construct(const std::vector<int>& _v) {  
    tag = VECTOR;  
    new (&v) std::vector<int>(_v);  
}
```

```
public:
```

```
template <typename T>  
svu(const T& x) {  
    construct(x);  
}
```

Introduction

Motivating
Example

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Generic
variant

Questions

```
public:
    template <typename R, typename F>
    R visit(F&& f) {
        switch (tag) {
            case STRING: return f(s);
            case VECTOR:  return f(v);
        }
    }
    // and a const version
```

- ▶ So the object you pass in has to have an **operator()** that takes a string and one that takes a vector<**int**>.
- ▶ Both operators must return something convertible to R.

Visitation example

```
struct Visitor {  
    int operator()(const string& s) { return s.size(); }  
    int operator()(const vector<int> v) { return v[0]; }  
};  
svu x("Hello");  
int i = x.visit<int>(Visitor()); // 5  
int i = x.visit<int>([](const auto& v) {  
    return v.size();  
});
```

Destruction

```
private:
    void destruct() {
        visit<void>([](auto&& x) {
            using T = std::decay_t<decltype(x)>;
            x.~T();
        });
    }
public:
    ~svu() { destruct(); }
```

- ▶ decay_t?
- ▶ This has no dependencies on the particular types we're using.

Assignment

```
svu& operator=(const svu& other) {  
    destruct();  
    other.visit<void>(  
        [this](auto&& v) { construct(v); });  
    return *this;  
}
```

Move semantics

- ▶ We can define move copy constructors and assignment operators
- ▶ And we should for efficiency's sake
- ▶ But it's tedious — just chuck some `&&s` and `std::moves` around the place

Inline visitors

- Recall how we defined a visitor earlier:

```
struct Visitor {  
    int operator()(const string& s) {return s.size();}  
    int operator()(const vector<int> v) {return v[0];}  
};
```

- Wouldn't it be nice to not have to explicitly define a struct for this?

```
x.visit<int>(  
    [] (const string& s) { return s.size(); },  
    [] (const vector<int> v) { return v[0]; });
```

- We can do this with a little recursive class template...

Overload set

- ▶ We start by declaring the class template. The template parameters are the types of the function-like objects that implement it:

```
template <typename... Fs>  
class overload_set;
```

- ▶ Then we define the base case for the recursion - an overload set with zero functions:

```
template <>  
class overload_set<> {  
public:  
    void operator()() = delete;  
};
```

Overload set

- Finally we define the inductive case — an overload set with $n + 1$ functions defined in terms of one with n :

```
template <typename F, typename... Fs>
class overload_set<F, Fs...>
    : private overload_set<Fs...>, private F {
public:
    explicit overload_set(F&& f, Fs&&... fs)
        : overload_set<Fs...>(std::forward<Fs>(fs)...),
          F(std::forward<F>(f)) {}

    using F::operator();
    using overload_set<Fs...>::operator();
};
```

- Now we can add an overload of our visit function to create an `overload_set` if we pass other than one argument.

```
template <typename R, typename... Fs>
R visit(Fs&&... fs) {
    return visit<R>(
        overload_set<Fs...>(std::forward<Fs>(fs)...));
}
// and a const version
```

Inline visitation

```
x.visit<int>(  
    [](const std::string& s) { return s.size(); },  
    [](const std::vector<int> v) { return v[0]; });
```

- Note how `s.size()` actually returns `size_t` but it gets implicitly converted to `int` — this is useful in many cases but consider `-Wconversion`.

- What if we want to visit two objects at once?

```
int plux(const svu& u1, const svu& u2) {  
    using namespace std;  
    return u1.visit<int>(  
        [&](const string& s1) {  
            return u2.visit<int>(  
                [&](const string& s2) { return s1.size() + s2.size(); },  
                [&](const vector<int>& v2) { return s1.size() + v2[0]; });  
            },  
        [&](const vector<int> v1) {  
            return u2.visit<int>(  
                [&](const string& s2) { return v1.size() + s2[0]; },  
                [&](const vector<int>& v2) { return v1[0] + v2[0]; });  
            });  
    }  
}
```

- That makes me want to claw my eyes out.

► I want to be able to write:

```
int plux2(const svu& u1, const svu& u2) {  
    using namespace std;  
    auto visitor = make_multivisitor<int>(  
        [] (const string& s1, const string& s2) { return s1.size() + s2.size(); },  
        [] (const string& s1, const vector<int>& v2) { return s1.size() + v2[0]; },  
        [] (const vector<int> v1, const string& s2) { return v1.size() + s2[0]; },  
        [] (const vector<int> v1, const vector<int>& v2) { return v1[0] + v2[0]; });  
    return visitor(u1, u2);  
}
```

► This is much nicer because

- it scales gracefully to any number of visitees,
- it allows wildcards in any position and
- it only makes we want to claw one eye out.

► But how do we make one?

- It starts out pretty simple...

```
template <typename R, typename F>
class multivisitor {
private:
    F m_f;

public:
    explicit multivisitor(F&& f) : m_f(f) {}

    template <typename... Vs>
    auto operator()(const Vs&... args) {
        return collect(std::tuple<>(), args...);
    }
}
```

- ▶ The tricky bit is that we need to accumulate the results of visiting each variant until we've visited them all.

```
private:
template <typename T>
auto collect(const T& t) {
    return apply(m_f, t);
}

template <typename T, typename V, typename... Vs>
auto collect(const T& t, const V& arg, const Vs&... args) {
    return arg.template visit<R>([&](auto v) {
        return this->collect(
            std::tuple_cat(t, std::make_tuple(v)), args...);
    });
}
};
```

[Introduction](#)[Motivating
Example](#)[Discriminated
Unions](#)Plain old `union`
`union` + tag[Inline visitors](#)[Multi-
visitation](#)[Generic
variant](#)[Questions](#)

- ▶ So, that's all well and good if you want a variant that can hold a string or a vector<int>.
- ▶ But we wanted two different variants — one for states and one for events.
- ▶ So let's copy and paste and change the names, job done.

Hell, no!

We're just getting warmed up.

Construction

- ▶ We need a construct function for each type in the variant.
- ▶ So we go to our old friend, the recursive class template.
- ▶ The base case, for a variant with zero types:

```
template <typename I, I N, typename... Ts>
struct variant_construct;

template <typename I, I N>
struct variant_construct<I, N> {
    static I construct();
};
```

Construction

- And the inductive $n + 1$ case defined in terms of the n case:

```
template <typename I, I N, typename T, typename... Ts>
struct variant_construct<I, N, T, Ts...>
    : private variant_construct<I, N + 1, Ts...> {
    using super = variant_construct<I, N + 1, Ts...>;

    static I construct(void* storage, const T& value) {
        new (storage) T(value);
        return N;
    }
    using super::construct;
};
```

Visitation

Type-safe
Unions in
C++14

Toby Allsopp

```
template <typename I, I N, typename... Ts>
struct variant_visit;

template <typename I, I N>
struct variant_visit<I, N> {
    template <typename R, typename F>
    static R visit_helper_const(I tag, const void*, F&&) {
        throw std::logic_error("variant tag invalid");
    }
};

template <typename I, I N, typename T, typename... Ts>
struct variant_visit<I, N, T, Ts...> : private variant_visit<I, N + 1, Ts...> {
    using super = variant_visit<I, N + 1, Ts...>;

    template <typename R, typename F>
    static R visit_helper_const(I tag, const void* storage, F&& f) {
        if (tag == N) {
            return f(*reinterpret_cast<const T*>(storage));
        }
        return super::template visit_helper_const<R>(tag, storage, std::forward<F>(f));
    }
};
```

Introduction

Motivating
Example

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Generic
variant

Questions

Storage

- ▶ I can't figure out how to use a `union` as storage.
- ▶ But we can use a very handy template:

```
typename std::aligned_union_t<0, Ts...> storage;
```

- ▶ The size needed for the tag depends on how many types are in the variant.

```
template <uintmax_t N, typename Enable = void>
struct smallest_unisnged_type;

template <uintmax_t N>
struct smallest_unisnged_type<
    N, typename std::enable_if_t<N <= std::numeric_limits<uint8_t>::max()>> {
    using type = uint8_t;
};

template <uintmax_t N>
using smallest_unisnged_type_t = typename smallest_unisnged_type<N>::type;
```

variant part 1

Type-safe
Unions in
C++14

Toby Allsopp

```
template <typename... Ts>
struct variant_helper {
    using tag_type = smallest_unisnged_type_t<sizeof...(Ts)>;
    using super_construct = variant_construct<tag_type, 0, Ts...>;
    using super_visit = variant_visit<tag_type, 0, Ts...>;
};

template <typename... Ts>
class variant : private detail::variant_helper<Ts...>::super_construct,
               private detail::variant_helper<Ts...>::super_visit {
    using helper = detail::variant_helper<Ts...>;
    using super_construct = typename helper::super_construct;
    using super_visit = typename helper::super_visit;

    typename std::aligned_union_t<0, char, Ts...> storage;
    typename helper::tag_type tag;

    using super_construct::construct;

    void destruct() {
        std::move(*this).template visit<void>([this](auto&& v) {
            using T = std::decay_t<decltype(v)>;
            reinterpret_cast<T*>(&storage)->~T();
        });
    }
}
```

Introduction

Motivating
Example

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Generic
variant

Questions

variant part 2

```
public:
    template <typename T, typename = decltype(construct(
        &storage, std::forward<T>(std::declval<T>())))>
    variant(T&& value) {
        tag = construct(&storage, std::forward<T>(value));
    }

    variant(const variant& other) {
        other.visit<void>([this](auto&& value) { tag = construct(&storage, value); });
    }

    variant& operator=(const variant& other) {
        destruct();
        other.visit<void>([this](auto&& value) { tag = construct(&storage, value); });
        return *this;
    }

    ~variant() { destruct(); }

    template <typename R, typename F>
    auto visit(F&& f) const& {
        return super_visit::template visit_helper_const<R>(tag, &storage,
            std::forward<F>(f));
    }
};
```

Variant type

- ▶ Finally, we can define the types for our robot's state and events

```
using state = variant<off, idle, turning>;  
using on = variant<idle, turning>;  
  
using event = variant<turn_on, turn_off, start_turning, reset,  
                     heading_changed>;
```

- ▶ That `on` type comes in handy for the state transitions.

State transitions

```
state transition(const state& s, const event& e) {  
    return make_multivisitor<state>(  
        [](off,      turn_on)      { return idle{}; },  
        [](off,      auto)         { return off{}; },  
        [](on,       turn_off)     { return off{}; },  
        [](auto s,   turn_on)      { return s; },  
        [](on,       reset)        { return idle{}; },  
        [](on,       start_turning e) { return turning{e.target}; },  
        [](idle s,   heading_changed) { return s; },  
        [](turning s, heading_changed e) -> state {  
            if (std::abs(e.heading - s.target) < .1f) {  
                return idle{};  
            } else {  
                return s;  
            }  
        }) (s, e);  
}
```

Introduction

Motivating
Example

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Generic
variant

Questions

Questions?

Appendix

Extra material
References

Appendix

Extra material
References

References

- ▶ Boost