

Type-safe Unions in C++

Toby Allsopp
toby@mi6.gen.nz

Auckland C++ Meetup
19 July 2016

This talk

Type-safe
Unions in C++

Toby Allsopp

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old `union`
`union + tag`

Inline visitors

Multi-
visitation

Questions

Overview

Introduction

Motivating Example

A Touch of Theory

Discriminated Unions

Plain old **union**
union + tag

Inline visitors

Multi-visitation

Questions

Type-safe
Unions in C++

Toby Allsopp

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old **union**
union + tag

Inline visitors

Multi-
visitation

Questions

A robot

► States

```
#include <iostream>
```

```
struct idle {};
```

```
struct turning {  
    float targetHeading;  
};
```

```
std::ostream& operator<<(std::ostream& os, const idle&) {  
    return os << "idle{}";  
}
```

```
std::ostream& operator<<(std::ostream& os, const turning& s) {  
    return os << "turning{" << s.targetHeading << "}";  
}
```

► Events

```
#include <iostream>
```

State transitions

Start state	Event	End state
Any	<code>start_turning{h}</code>	<code>turning{h}</code>
Any	<code>stop_turning</code>	<code>idle</code>
<code>turning{h}</code>	<code>heading_changed{h}</code>	<code>idle</code>
<code>s</code>	<code>heading_changed</code>	<code>s</code>

A **type** is EITHER

- ▶ a set of n primitive values $T = \{v_1, v_2, \dots, v_n\}$,
- ▶ a **product** of n other types $T = T_1 \times T_2 \times \dots \times T_n$, or
- ▶ a **sum** of n other types $T = T_1 \cup T_2 \cup \dots \cup T_n$.

Sum types

- ▶ To represent the state of the robot, we want a type whose values are **either** `idle` **or** `turning`.
- ▶ This is called a **sum type** because its cardinality is the sum of the cardinalities of its constituent types.
- ▶ We could also call it a **union type** because the set of its possible values is the union of the sets of values of its constituent types.

Indiscriminate **union**

Type-safe
Unions in C++

Toby Allsopp

```
union u {  
    int i;  
    double d;  
}  
u x;  
x.i = 3;  
double d = x.d; // UNDEFINED BEHAVIOUR
```

- ▶ Can't tell what was last stored but you'd better know!
- ▶ Can be used for bit twiddling on specific platforms
- ▶ No guarantees from the standard

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old **union**
union + tag

Inline visitors

Multi-
visitation

Questions

More **union** gotchas

Type-safe
Unions in C++

Toby Allsopp

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old **union**
union + tag

Inline visitors

Multi-
visitation

Questions

```
union u {  
    std::string s;  
    std::vector<int> v;  
} x;  
x.s = "Hello"; // KABOOM!  
new (&x.s) std::string("Hello"); // OK  
x.s = "Goodbye"; // OK  
u y = x; // NOPE - copy constructor is deleted  
new (&x.v) std::vector<int>{1, 2, 3}; // LEAK!  
x.s.~std::string(); // DO THIS FIRST
```

union + class + enum

Type-safe
Unions in C++

Toby Allsopp

```
class svu {  
    private:  
        enum { STRING, VECTOR } tag;  
        union {  
            std::string s;  
            std::vector<int> v;  
        };  
};
```

- ▶ Because it's so unsafe, let's bundle it up into a class and only expose safe operations.
- ▶ We need to keep track of which union member we have initialized - we call this a **tag**.

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old **union**
union + tag

Inline visitors

Multi-
visitation

Questions

Construction

```
private:
```

```
void construct(const std::string& _s) {  
    tag = STRING;  
    new (&s) std::string(_s);  
}  
void construct(const std::vector<int>& _v) {  
    tag = VECTOR;  
    new (&v) std::vector<int>(_v);  
}
```

```
public:
```

```
template <typename T>  
svu(const T& x) {  
    construct(x);  
}
```

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Questions

Visitation

```
public:
    template <typename R, typename F>
    R visit(F&& f) {
        switch (tag) {
            case STRING: return f(s);
            case VECTOR:  return f(v);
        }
    }
    // and a const version
```

- ▶ So the object you pass in has to have an **operator()** that takes a string and one that takes a vector<**int**>.
- ▶ Both operators must return something convertible to R.

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Questions

Visitation example

```
struct Visitor {  
    int operator()(const string& s) { return s.size(); }  
    int operator()(const vector<int> v) { return v[0]; }  
};  
svu x("Hello");  
int i = x.visit<int>(Visitor()); // 5  
int i = x.visit<int>([](const auto& v) {  
    return v.size();  
});
```

Destruction

```
private:
    void destruct() {
        visit<void>([](auto&& x) {
            using T = std::decay_t<decltype(x)>;
            x.~T();
        });
    }
public:
    ~svu() { destruct(); }
```

- ▶ decay_t?
- ▶ This has no dependencies on the particular types we're using.

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old union
union + tag

Inline visitors

Multi-
visitation

Questions

Assignment

Type-safe
Unions in C++

Toby Allsopp

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old `union`
`union + tag`

Inline visitors

Multi-
visitation

Questions

```
svu& operator=(const svu& other) {  
    destruct();  
    other.visit<void>(  
        [this](auto&& v) { construct(v); });  
    return *this;  
}
```

Move semantics

Type-safe
Unions in C++

Toby Allsopp

Introduction

Motivating
Example

A Touch of
Theory

Discriminated
Unions

Plain old `union`

`union + tag`

Inline visitors

Multi-
visitation

Questions

Inline visitors

- Recall how we defined a visitor earlier:

```
struct Visitor {  
    int operator()(const string& s) {return s.size();}  
    int operator()(const vector<int> v) {return v[0];}  
};
```

- Wouldn't it be nice to not have to explicitly define a struct for this?

```
x.visit<int>(  
    [] (const string& s) { return s.size(); },  
    [] (const vector<int> v) { return v[0]; });
```

- We can do this with a little recursive class template...

Overload set

- ▶ We start by declaring the class template. The template parameters are the types of the function-like objects that implement it:

```
template <typename... Fs>  
class overload_set;
```

- ▶ Then we define the base case for the recursion - an overload set with zero functions:

```
template <>  
class overload_set<> {  
public:  
    void operator()() = delete;  
};
```

Overload set

- Finally we define the inductive case — an overload set with $n + 1$ functions defined in terms of one with n :

```
template <typename F, typename... Fs>
class overload_set<F, Fs...>
    : private overload_set<Fs...>, private F {
public:
    explicit overload_set(F&& f, Fs&&... fs)
        : overload_set<Fs...>(std::forward<Fs>(fs)...),
          F(std::forward<F>(f)) {}

    using F::operator();
    using overload_set<Fs...>::operator();
};
```

- ▶ Now we can add an overload of our visit function to create an `overload_set` if we pass more than one argument.

```
template <
    typename R, typename... Fs,
    typename = std::enable_if_t<sizeof...(Fs) >= 2>>
R visit(Fs&&... fs) {
    return visit<R>(
        overload_set<Fs...>(std::forward<Fs>(fs)...));
}
// and a const version
```

- ▶ In fact, we could make the original version private and use this one unconditionally — why not?

Inline visitation

```
x.visit<int>(
    [](const std::string& s) { return s.size(); },
    [](const std::vector<int> v) { return v[0]; });
```

- Note how `s.size()` actually returns `size_t` but it gets implicitly converted to `int` — this is useful in many cases but consider `-Wconversion`.

- What if we want to visit two objects at once?

```
int plux(const svu& u1, const svu& u2) {  
    using namespace std;  
    return u1.visit<int>(  
        [&](const string& s1) {  
            return u2.visit<int>(  
                [&](const string& s2) { return s1.size() + s2.size(); },  
                [&](const vector<int>& v2) { return s1.size() + v2[0]; });  
            },  
        [&](const vector<int> v1) {  
            return u2.visit<int>(  
                [&](const string& s2) { return v1.size() + s2[0]; },  
                [&](const vector<int>& v2) { return v1[0] + v2[0]; });  
            });  
    }  
}
```

- That makes me want to claw my eyes out.

► I want to be able to write:

```
int plux2(const svu& u1, const svu& u2) {  
    using namespace std;  
    auto visitor = make_visitor<int>(  
        [](const string& s1,      const string& s2)      { return s1.size() + s2.size(); },  
        [](const string& s1,      const vector<int>& v2) { return s1.size() + v2[0]; },  
        [](const vector<int> v1,   const string& s2)      { return v1.size() + s2[0]; },  
        [](const vector<int> v1,   const vector<int>& v2) { return v1[0] + v2[0]; });  
    return visitor(u1, u2);  
}
```

- This is much nicer because
- it scales gracefully to any number of visitees,
 - it allows wildcards in any position and
 - it only makes we want to claw one eye out.
- But how do we make one?

```
template <typename R, typename F>
class visitor {
private:
    F m_f;

public:
    explicit visitor(F&& f) : m_f(f) {}

    template <typename... Vs>
    auto operator()(const Vs&... args) {
        return collect(std::tuple<>(), args...);
    }
}
```



```
private:
    template <typename T>
    auto collect(const T& t) {
        return apply(m_f, t);
    }

    template <typename T, typename V, typename... Vs>
    auto collect(const T& t, const V& arg,
                const Vs&... args) {
        return arg.template visit<R>([&](auto v) {
            return this->collect(
                std::tuple_cat(t, std::make_tuple(v)),
                args...);
        });
    }
};
```

```
template <typename Callable, typename Tuple,  
          size_t... I>  
auto apply_impl(Callable&& f, Tuple&& t,  
                std::index_sequence<I...>) {  
    return f(std::get<I>(t)...);  
}  
  
template <typename Callable, typename Tuple>  
auto apply(Callable&& f, Tuple&& t) {  
    using is = std::make_index_sequence<  
        std::tuple_size<std::decay_t<Tuple>>::value>;  
    return apply_impl(std::forward<Callable>(f),  
                      std::forward<Tuple>(t), is());  
}
```

```
state transition(const state& s, const event& e) {  
    return make_visitor<state>(make_overload_set(  
        [](auto,      start_turning e) {  
            return turning{e.target};  
        },  
        [](auto,      stop_turning) { return idle{}; },  
        [](auto s,     heading_changed) { return s; },  
        [](turning s, heading_changed e) -> state {  
            if (std::abs(e.heading - s.target) < .1f) {  
                return idle{};  
            } else {  
                return s;  
            }  
        }  
    ))(s, e);  
}
```

Questions?

Appendix

Extra material
References

References

- ▶ Boost